

AI Act Transparency Requirements

Transparency obligations for AI systems that interact with people: chatbots, synthetic content, deepfakes, emotion recognition and biometric categorisation. Article 50 of Regulation (EU) 2024/1689 (Article 52 of the 2021 proposal). Version of 29 September 2026.

Use of this document. Internal NormaAI summary for consultants, not legal advice. Check statements against the official text on EUR-Lex (CELEX 32024R1689) before advising a client.

Article numbering. The transparency rules are in **Article 50** of the adopted AI Act. In the Commission's 2021 proposal and in many early articles and training materials they appeared as **Article 52**. In the final Regulation (EU) 2024/1689, Article 52 is a different provision: the procedure for designating general-purpose AI models with systemic risk. When a client cites "Article 52 transparency", they mean Article 50 of the final text.

1. Overview

Article 50 sets transparency duties for four kinds of AI use where people could be misled about whether they deal with a machine, or whether content is real. These systems are sometimes called **limited risk**. Article 50 is not limited to systems outside the high-risk category: it applies in addition to the high-risk rules when a system is both. The duties are placed on providers or on deployers depending on the situation.

Situation	Who must act	What must be done	Article
Chatbots and other AI systems interacting directly with people	Provider	Design the system so that people are informed they are interacting with an AI system, unless obvious.	50(1)
Synthetic content: AI-generated audio, image, video or text	Provider	Mark outputs in a machine-readable format and make them detectable as artificially generated or manipulated.	50(2)
Emotion recognition and biometric categorisation systems	Deployer	Inform the people exposed to the system of its operation and process personal data in line with data protection law.	50(3)
Deepfakes	Deployer	Disclose that the content has been artificially generated or manipulated.	50(4), first subparagraph
AI-generated text published on matters of public interest	Deployer	Disclose that the text has been artificially generated or manipulated, unless there is human review or editorial control.	50(4), second subparagraph

2. Timing

Article 50 applies from **August 2026**, the general date of application of the Act (Article 113). From that date the following apply: chatbot disclosure (Article 50(1)), machine-readable marking of synthetic content (Article 50(2)), the emotion recognition and biometric categorisation notice (Article 50(3)), and the deepfake and public-interest text disclosure (Article 50(4)). Fines for breaching Article 50 can be imposed under Article 99(4). Verify the exact dates on EUR-Lex before advising a client.

3. Chatbots and systems interacting with people (Article 50(1))

Rule. Providers must ensure that AI systems intended to **interact directly with natural persons** are designed and developed in such a way that the natural persons concerned are informed that they are interacting with an AI system. The duty is on the provider, so it must be built into the product.

Exception: obvious from context. The obligation does not apply where this is obvious from the point of view of a natural person who is **reasonably well-informed, observant and circumspect**, taking into account the circumstances and the context of use. A humanoid robot in a shop window with a clear label, or a well-known virtual game character, may qualify. A chatbot with a human name and a photo, answering on a bank website, does not.

Exception: law enforcement. The obligation does not apply to AI systems authorised by law to detect, prevent, investigate or prosecute criminal offences, subject to appropriate safeguards for third-party rights, unless those systems are available for the public to report a criminal offence.

3.1 Practical guidance for chatbot clients

- Show a **clear notice at the start of the conversation**, for example "You are chatting with an AI assistant". Article 50(5) requires the information to be given at the latest at the time of the first interaction or exposure.
- Make the notice **clear and distinguishable** and accessible, including for people with disabilities (Article 50(5)).
- Voice assistants and call bots need a spoken notice, not only a visual one.
- If a conversation moves from AI to a human agent and back, make it clear who is speaking.
- Do not design personas that deny being an AI. Beyond Article 50, deceptive techniques may fall under the prohibition in Article 5(1)(a).
- Keep evidence of the notice (screenshots, configuration records). It helps demonstrate compliance to authorities.

4. Synthetic content and machine-readable marking (Article 50(2))

Rule. Providers of AI systems, **including general-purpose AI systems**, that generate **synthetic audio, image, video or text content** must ensure that the outputs of the system are **marked in a machine-readable format and detectable as artificially generated or manipulated**.

The technical solutions must be **effective, interoperable, robust and reliable** as far as this is technically feasible, taking into account the specificities and limitations of different content types, the costs of implementation and the generally acknowledged state of the art, as may be reflected in relevant technical standards. In practice, this means techniques such as watermarks, metadata identifications, cryptographic methods for proving provenance, logging methods or fingerprints, or a combination of them.

4.1 Exceptions to Article 50(2)

- AI systems that perform an **assistive function for standard editing**, for example spell-checking or basic image adjustments.
- AI systems that **do not substantially alter** the input data provided by the deployer or its semantics.
- AI systems authorised by law to detect, prevent, investigate or prosecute criminal offences.

Who is affected. The duty falls on the provider of the generative system, not on the person who uses it. A company that builds a product on top of a general-purpose model and offers image, audio, video or text

generation to customers is a provider for the purposes of this Article and must make sure the marking is present in the outputs of its product. The Commission has been encouraged to support the drawing up of a **code of practice on the detection and labelling of artificially generated content** (Article 50(7)) and may approve it by implementing act. Consultants should check the current status of this code before advising on technical approaches.

5. Emotion recognition and biometric categorisation (Article 50(3))

Rule. Deployers of an **emotion recognition system** or a **biometric categorisation system** must inform the natural persons exposed thereto of the operation of the system. They must also process personal data in accordance with Regulations (EU) 2016/679 (GDPR) and (EU) 2018/1725 and Directive (EU) 2016/680, as applicable.

Term	Definition	Article 3
Emotion recognition system	An AI system for the purpose of identifying or inferring emotions or intentions of natural persons on the basis of their biometric data.	Art. 3(39)
Biometric categorisation system	An AI system for the purpose of assigning natural persons to specific categories on the basis of their biometric data, unless ancillary to another commercial service and strictly necessary for objective technical reasons.	Art. 3(40)
Biometric data	Personal data resulting from specific technical processing relating to the physical, physiological or behavioural characteristics of a natural person, such as facial images or dactyloscopic data.	Art. 3(34)

Exception. The duty does not apply to systems used for biometric categorisation and emotion recognition that are permitted by law to detect, prevent or investigate criminal offences, subject to safeguards and in accordance with Union law.

Important: transparency is not enough on its own. These systems face stricter rules elsewhere in the Act:

- **Emotion recognition in the workplace and in education institutions is prohibited**, except for medical or safety reasons (Article 5(1)(f)). A notice cannot make it lawful.
- **Biometric categorisation inferring race, political opinions, trade union membership, religious or philosophical beliefs, sex life or sexual orientation is prohibited** (Article 5(1)(g)).
- Outside those prohibitions, emotion recognition and biometric categorisation systems are **high-risk** under Annex III, point 1(b) and (c), and must meet the requirements of Articles 8 to 15 from August 2026, in addition to the Article 50(3) notice.
- Use of such a system in retail, marketing or entertainment (for example measuring customer reactions to an advertisement) is therefore typically high-risk and needs the notice to the people exposed.

6. Deepfakes and AI-generated text (Article 50(4))

Deepfake definition (Article 3(60)). AI-generated or manipulated image, audio or video content that resembles existing persons, objects, places, entities or events and would falsely appear to a person to be authentic or truthful.

Rule for deepfakes. Deployers of an AI system that generates or manipulates image, audio or video content constituting a deepfake must **disclose that the content has been artificially generated or manipulated**.

Artistic, creative, satirical or fictional works. Where the content forms part of an evidently artistic, creative, satirical, fictional or analogous work or programme, the disclosure duty is limited to disclosing the existence of such generated or manipulated content **in an appropriate manner that does not hamper the display or enjoyment of the work.**

Rule for AI-generated text. Deployers of an AI system that generates or manipulates **text published with the purpose of informing the public on matters of public interest** must disclose that the text has been artificially generated or manipulated.

- **Editorial exemption:** the text disclosure does not apply where the AI-generated content has undergone a process of **human review or editorial control** and where a natural or legal person holds **editorial responsibility** for the publication of the content.
- **Law enforcement exemption:** systems authorised by law to detect, prevent, investigate or prosecute criminal offences are exempt from the deepfake duty.
- The information must be given at the latest at the time of the first interaction or exposure, in a clear and distinguishable manner, and meet applicable accessibility requirements (Article 50(5)).

6.1 Who does what: provider or deployer

Actor	Duty	Example
Provider of a generative system	Embed machine-readable marking in outputs (Art. 50(2)).	An image generation tool adds a watermark or provenance metadata to every image.
Deployer publishing a deepfake	Disclose visibly or audibly that content is artificial (Art. 50(4)).	A marketing team publishes a synthetic video of a spokesperson and labels it as AI-generated.
Deployer publishing AI text on public interest matters	Disclose AI generation unless human editorial control (Art. 50(4)).	A news site publishes an automatically generated market report without editorial review, with a disclosure.

7. General rules for all Article 50 duties

- **Timing and form (Article 50(5)):** information must be provided to the persons concerned in a clear and distinguishable manner at the latest at the time of the first interaction or exposure, and must conform to accessibility requirements.
- **Relation to other rules (Article 50(6)):** Article 50 does not affect the requirements and obligations in Chapter III (high-risk systems), nor other transparency obligations in Union or national law.
- **Codes of practice (Article 50(7)):** the AI Office encourages and facilitates codes of practice at Union level to facilitate the effective implementation of the obligations on detection and labelling of artificially generated content. The Commission may assess codes and approve them by implementing act.
- **Guidelines (Article 96):** the Commission may issue guidelines on the practical implementation of Article 50.
- **Open source (Article 2(12)):** the open-source exemption from the Act does not apply to AI systems subject to Article 50.

8. Related transparency duties outside Article 50

Provision	Content	Difference from Article 50
Art. 13	Transparency and information to deployers for high-risk systems: instructions for use.	Addresses deployers, not the people affected. See AI_Act_High_Risk_Obligations .
Art. 26(11)	Deployers of high-risk systems that take or assist decisions about natural persons must inform them that they are subject to the use of such a system.	Applies only to high-risk systems.
Art. 86	Right of an affected person to a clear and meaningful explanation of the role of a high-risk system in a decision.	An individual right after a decision, not a general notice.
Art. 53	Providers of general-purpose AI models must keep technical documentation, give information to downstream providers, put in place a copyright policy and publish a sufficiently detailed summary of the content used for training .	Model-level transparency, in Chapter V.
Art. 55	Additional duties for providers of general-purpose AI models with systemic risk : model evaluation, adversarial testing, serious incident reporting, cybersecurity.	Chapter V. The systemic risk designation procedure is in Article 52 of the final text.

9. Compliance checklist for a client

#	Question	If yes, action
1	Does our product talk or write to people using AI (chatbot, voice bot, virtual agent)?	Add a clear AI notice at first interaction. Document it (Art. 50(1)).
2	Does our product generate audio, images, video or text for users?	Implement machine-readable marking. Article 50(2).
3	Do we publish deepfakes, or AI text on public interest matters?	Add visible disclosure. Set up editorial review for text if we want to rely on the exemption (Art. 50(4)).
4	Do we use emotion recognition or biometric categorisation?	First check that it is not prohibited (Art. 5(1)(f) and (g)). If allowed, inform exposed persons (Art. 50(3)), handle personal data lawfully and plan for high-risk duties.
5	Is the system also in Annex III?	Both regimes apply. See AI_Act_Risk_Classification .
6	Do we resell or integrate a third-party generative model?	Check with the model provider who provides the marking, and confirm we do not remove it. We may be a provider under Art. 50(2).

10. Consequences of non-compliance

Breach of the transparency obligations for providers and deployers under Article 50 is subject to administrative fines of up to **EUR 15 million or 3% of total worldwide annual turnover** for the preceding financial year, whichever is higher (Article 99(4)(g)). For SMEs and start-ups, the fine is the lower of the two amounts (Article 99(6)). Market surveillance authorities enforce Article 50. The full penalty structure is in [AI_Act_Penalties_Enforcement](#).

11. Frequently asked questions

Do we have to tell customers our support chatbot is AI?

Yes, unless it is obvious to a reasonably well-informed, observant and circumspect person (Article 50(1)). Most website chatbots need a clear notice. The duty applies from 2 August 2026.

Do we need to label an email that AI helped us draft?

Article 50 does not require a label on every AI-assisted text. Text disclosure applies to AI-generated text published to inform the public on matters of public interest, and even then it is not needed where a human reviews it and someone holds editorial responsibility (Article 50(4)). Other rules, such as consumer law or the GDPR, may still apply.

Are we responsible for watermarking images our users generate with our app?

Yes, if you are the provider of the generative system. Article 50(2) requires the outputs to be marked in a machine-readable format.

Does a deepfake label make a harmful deepfake lawful?

No. Article 50(4) only sets a disclosure duty. Content can still be unlawful under Article 5 (for example the new ban on non-consensual intimate imagery), under the GDPR, or under national criminal, defamation and image rights law.

Where does the AI Act say "Article 52" for transparency?

It does not, in the final text. The 2021 proposal put transparency in Article 52. The adopted Regulation (EU) 2024/1689 moved it to Article 50. Article 52 of the adopted text concerns the procedure for general-purpose AI models with systemic risk.

12. Sources

- Regulation (EU) 2024/1689, Articles 3(34), 3(39), 3(40), 3(60), 5, 50, 52, 53, 55, 96, 99, EUR-Lex, CELEX 32024R1689.
- European Commission, AI Office, code of practice work on marking and labelling of AI-generated content, digital-strategy.ec.europa.eu.